

Emulation Framework for AIoT Federated Learning

Brendan Ang Wei Jie

January 22, 2024

1 Abstract

Federated learning allows a fleet of devices to collaborate towards a globally trained machine learning model. Research has continued to produce novel federated learning algorithms to tackle different issues in FL such as heterogeneity and learning over data from non-identical distributions. Performance of these algorithms depend in part on the system parameters used in FL such as number of clients and number of passes. Furthermore, a realistic benchmark would require one to procure a large fleet of devices. This work seeks to introduce a software emulation framework to streamline the process of building a fleet of clients and allow easy testing of FL system parameters for configuration of optimal values.

2 Acknowledgements

Contents

1	Abstract	2
2	Acknowledgements	3
3	Introduction	5
3.1	Problem	5
3.2	Literary Review	6
3.3	QEMU	6
3.4	Zephyr OS	7
4	Implementation	7
4.1	Client	11
4.2	Server	11
5	Experiments	12
6	Limitations	12
6.1	Long training times	12

List of Figures

1	Host-guest emulation architecture	10
2	Client initialization sequence	11
3	HTTP endpoints	12

3 Introduction

Federated learning emerged as a method for solving key issues with the standard centralized learning approach. Some of these issues are (1) Preserving user data privacy: a centralized training approach involves the need for the central machine performing the computation to have full access to all the data. With FL, data never leaves each individual client’s device. Instead, only the updated weights are shared to form the global model.(2) Scalability: FL enables leveraging a network to perform computation in parallel. The distributed nature of FL poses a few key challenges. Some of these include:

- Limited computing resources: individual devices may have constraints on memory, limiting the size of the local model it is able to train. Constraints on computing power can also lead to longer training times. This is particularly so for Internet of Things (IoT) devices such as sensors, where their embedded nature leads to a limited size and power.
- Network limitations: communication speed can become the bottleneck for performance as IoT devices rely on unstable wireless communication networks. Furthermore, data constraints can exist such that the number of bytes sent may be a cost inducing factor.

Hence, development of FL algorithms often seek to advance progress towards solving these aforementioned issues.

3.1 Problem

However, prior to effectively deploying FL, different factors including the convergence rate and model accuracy needs to be well studied, giving rise to the need for simulation. In addition, the simulation of FL in the distributed setting involves dealing with issues which do not arise in datacenter ML research. These include running on different simulated devices each with potentially varying amount of data. Furthermore, metrics such as number of bytes uploaded and downloaded by the device as well as the ability to simulate real-world issues such as client drop-out can also be important for proposed FL algorithms to handle.

3.2 Literary Review

FedML[?] is a research-oriented library which supports various algorithms and 3 platforms, on-device training for IoT and mobile devices, distributed computing and single-machine simulation. On-device training is supported only on Raspberry Pi 4 and NVIDIA Jetson Nano which limits hardware validation to these 2 devices. Simulation is offered through the FedML Parrot[?], which offers an accelerated simulation framework employing multiple optimizations to improve simulation speed.

FLSim[?] aims to provide a simulation framework for FL by offering users a set of software components which can then be mixed and matched to create a simulator for their use case. Developers need only define the data, model and metrics reported, and FL system parameters can be altered in a JSON configuration file. FLUTE[?] adds additional features by allowing users to gain access to cloud based compute and data. However, these simulation solutions are consequently hardware-agnostic. Without taking into account the specific platforms which FL is performed on, these simulators cannot provide insight into whether the system will work in the production environment.

Here, this work proposes a new emulation framework zfl, which aims to achieve the following properties:

- Provide better insight into platform specific support by running FL on emulated hardware.
- Collect useful metrics during the FL lifecycle.
- Support implementation of real word issues such as variable data and client drop-off.

3.3 QEMU

QEMU[?] is an open source machine emulator and virtualizer. It enables system emulation, where it provides a virtual model of an entire machine (CPU, memory and emulated devices) to run a guest OS. In this mode the CPU may be fully emulated, or it may work with a hypervisor to allow the guest to run directly on the host CPU. In zfl, QEMU with full CPU emulation is used without a hypervisor as part of its software architecture to emulate hardware.

3.4 Zephyr OS

To emulate the issue of limited computing resources, it is important to make use of system runtimes used by those devices. In particular, the type of operating system used will help to ensure that the kernel is lightweight and configurable. Zephyr OS[?] is one such OS. It is based on a small-footprint kernel designed for use on resource-constrained and embedded systems: from simple embedded environmental sensors and LED wearables to sophisticated embedded controllers, smart watches, and IoT wireless applications. Furthermore, Zephyr is highly configurable, allowing the user to choose only the specific kernel services required, and also delve into lower level memory allocations of the system RAM. In zfl, client code will be written to work in the Zephyr OS environment.

4 Implementation

zfl aims to provide the ability to emulate the traditional FedAvg[?] algorithm, with multiple clients communicating with a central server which performs the aggregation. To run on Zephyr OS, the entire software stack is developed in the C programming language. This also makes it suitable to run on embedded devices.

Next the framework needed a method for spawning an arbitrary number of QEMU instances, and allowing these instances to communicate back to the server. When called in client mode, zfl accomplishes this by making use of the ‘fork’ and ‘exec’ pattern with the desired number of clients.

However, each instance of QEMU persists in an isolated network different from the host PC, and is not able to communicate. We can bypass this limitation using a network bridge to act as a virtual network device forwarding packets between connected network devices. The network bridge is set up on the host under the name zfl with a set of network parameters using the ip command line utility.

```
1 ip link add $INTERFACE type bridge
2 ip addr add $IPV4_ADDR_1 dev $INTERFACE
3 ip link set enp61s0 master $INTERFACE
4 ip route add $IPV4_ROUTE_1 dev $INTERFACE > /dev/null 2>&1
```

```
5 ip link set dev $INTERFACE up
```

Listing 1: Network bridge setup

To tell QEMU to use it, we pass the name of the bridge along with a randomly generated MAC address as arguments to the `-nic` flag

```
-nic bridge,model=e1000,mac=%s,br=zfl
```

Although each QEMU client is now able to communicate with the host via the nic adapter, we still needed a way to monitor the output of each instance. One method is to transmit output and logs over the network. However, this would not allow important crash logs and stacktrace information to be transmitted as the application software would have shutdown. To overcome this, the host creates a named or FIFO pipe [?] for each client which is passed in to QEMU through the `-serial` flag. Now, all standard output goes through the named pipe, ready to be read by the host. Another outcome of this is that the host is now also able to send input to each instance, whose importance will be described in the next section.

```
1 pid_t child = fork();
2 if (child < 0) {
3     printf("ERROR: could not fork client %d: %s\n", i,
4           strerror(errno));
5     return 1;
6 }
7 ...
8 // generate serial arguments
9 char serial_arg[80];
10 snprintf(serial_arg, 80, "pipe:%s", pipe_path);
11
12 // generate nic arguments
13 char nic_arg[100];
14 char *mac = generate_random_mac();
15 snprintf(nic_arg, sizeof(nic_arg), "bridge,model=e1000,mac=%s
16         ,br=zfl", mac);
17
18 // start client as new process
19 execlp("qemu-system-i386", "qemu-system-i386",
20       "-m", "15", "-cpu", "qemu32,+nx,+pae", "-machine", "
21       q35",
22       "-device", "isa-debug-exit,iobase=0xf4,iosize=0x04",
```



```

23     "-no-reboot", "-nographic", "-no-acpi",
24
25     "-serial", serial_arg,
26
27     "-nic", nic_arg,
28
29     "-kernel", "./zflclient/out/zephyr/zephyr.elf",
30
31     NULL);

```

Listing 2: Client forking process

In addition to the MAC address configuration, each client needed a unique Internet Protocol Version 4 (IPv4) address in order to establish TCP based connections with the central server. Samples provided by Zephyr OS describe a way to achieve this by setting a compile time configuration flag `CONFIG_NET_CONFIG_MY_IPV4_ADDR`. However, this is impractical to scale to a number of clients, since each client would need a separately compiled binary. Instead, we can assign the IPv4 address dynamically using the built-in Zephyr network function:

`net_if_ipv4_addr_add`

To achieve this, each client starts off as a Zephyr shell instance and a user-defined command is registered as a way to start the main program. The desired IP address is obtained using the command line argument in listing 3.

```

1 SHELL_CMD_ARG_REGISTER(run, NULL, "Run with IPv4 address",
   run, 4, 0);

```

Listing 3: Registering user defined command "run" to the function pointer run

The complete architecture is illustrated in figure 1.

```

1     char *addr_str = argv[1];
2     LOG_INF("instance ipaddr is %s", addr_str);
3     struct in_addr addr;
4     zsock_inet_pton(AF_INET, addr_str, &addr);
5     if (!net_if_ipv4_addr_add(net_if_get_default(), &addr,
   NET_ADDR_MANUAL, UINT32_MAX)) {
6         LOG_ERR("failed to add %s to interface", addr_str);
7         return -1;
8     }

```

Listing 4: the "run" function which performs IP address assignment at runtime




Figure 1: Host-guest emulation architecture




Figure 2: Client initialization sequence

4.1 Client

Each client establishes TCP socket connection to the central server and obtains their assigned ID and training data according to that ID. Once done, each client starts a HTTP server and marks itself as ready to begin the training round. The complete initialization process is illustrated in figure 2.

4.2 Server

The central server runs on the host machine without needing additional configuration. Its primary purpose is to serve as a HTTP server. Implemented endpoints are described in figure 3. After assigning each connecting client their ID and training data, it performs the function `start_round` at an interval of 10 seconds. Each time, the server pings all previously connected clients to check if they are ready to start the training round. When enough clients are ready, it sends a HTTP post request to the `train` endpoint of the client, which triggers the training function. Once the client has completed training, it sends the resulting weights of the local model to the HTTP POST endpoint `results` of the server, who then performs the aggregation. The sequence of events for a single round is illustrated in figure ??.

Client		
Endpoint	HTTP Method	Description
/start	POST	Initiates training round using updated weights with JSON body: {weights: int}
/ready	GET	Gets client ready status
Server		
/results?id=	POST	Client id submits results with JSON body: {round: int, weights: string}
/id	GET	Obtain an id for round participation
/training-data?id=	GET	Obtain training data for id
/training-labels?id=	GET	Obtain training labels for id

Figure 3: HTTP endpoints

5 Experiments

The framework is tested using FedAvg over the MNIST digit recognition data set. Due to memory constraints, clients are configured with a neural network architecture with 1 hidden layer of 16 nodes. The underlying neural network implementation utilizes nn.h[?], an open source educational neural network C library. Additionally, implementations for the softmax activation and cross entropy were made to make it suitable for machine learning with MNIST categorical data.

6 Limitations

6.1 Long training times

Unoptimized client training implementations has resulted in long training times required per FL round. This is because we were unable to find standard implementations for machine learning algorithms that can be easily ported for

use in the Zephyr OS. Zephyr OS provides external modules for supporting Tensorflow Lite Micro[?] and has examples for running pre-trained neural network on their platform. However, on-device training is not supported[?].